

```
import numpy as np

num_rows = 4
num_cols = 4
num_states = num_rows * num_cols # 4x4 grid
num_actions = 4 # Up, Down, Left, Right
num_episodes = 1000
alpha = 0.6 # Learning rate
gamma = 0.9 # Discount factor
epsilon = 0.1 # Epsilon for epsilon-greedy policy
```

```
# Initialize Q-values for SARSA and Q-learning
Q_sarsa = np.zeros((num_states, num_actions), dtype=float)
Q_qlearning = np.zeros((num_states, num_actions), dtype=float)

# Define rewards for each state
rewards = np.random.randint(1, 10, num_states)

# Define the goal state and obstacles
goal_state = num_states - 1
obstacles = [5, 7, 10]

# Update rewards for goal state and obstacles
rewards[goal_state] = 10
for obstacle in obstacles:
    rewards[obstacle] = -10
```

```
def epsilon_greedy_policy(state, Q_values, epsilon):
    if np.random.rand() < epsilon:
        return np.random.randint(0, num_actions)
    else:
        return np.argmax(Q_values[state])

def sarsa(num_states, num_actions, num_episodes, alpha, gamma, epsilon):
    for episode in range(num_episodes):
        state = np.random.randint(0, num_states)
        action = epsilon_greedy_policy(state, Q_sarsa, epsilon)

        while True:
            next_state = (state + action) % num_states
            next_action = epsilon_greedy_policy(next_state, Q_sarsa, epsilon)

            reward = rewards[next_state]

            Q_sarsa[state, action] = (1 - alpha) * Q_sarsa[state, action] + \
                alpha * (reward + gamma * Q_sarsa[next_state, next_action])

            state = next_state
            action = next_action

            if state == num_states - 1:
                break

        return Q_sarsa

def q_learning(num_states, num_actions, num_episodes, alpha, gamma, epsilon):
    for episode in range(num_episodes):
        state = np.random.randint(0, num_states)

        while True:
            action = epsilon_greedy_policy(state, Q_qlearning, epsilon)

            next_state = (state + action) % num_states

            reward = rewards[next_state]

            Q_qlearning[state, action] = (1 - alpha) * Q_qlearning[state, action] + \
                alpha * (reward + gamma * np.max(Q_qlearning[next_state, :]))

            state = next_state

            if state == num_states - 1:
                break

        return Q_qlearning
```

```

learned_Q_values_sarsa = sarsa(num_states, num_actions, num_episodes, alpha, gamma, epsilon)
learned_Q_values_qlearning = q_learning(num_states, num_actions, num_episodes, alpha, gamma, epsilon)

print("Learned Q-values (SARSA):")
print(learned_Q_values_sarsa)
print("\nLearned Q-values (Q-learning):")
print(learned_Q_values_qlearning)

final_policy_sarsa = np.argmax(learned_Q_values_sarsa, axis=1)
final_policy_qlearning = np.argmax(learned_Q_values_qlearning, axis=1)

print("\nFinal Policy (SARSA):")
print(final_policy_sarsa)
print("\nFinal Policy (Q-learning):")
print(final_policy_qlearning)

```

```

Learned Q-values (SARSA):
[[ 54.42142128  55.06119396  56.75490549  56.87139762]
 [ 52.90366432  57.06915499  59.19421384  46.91883665]
 [ 47.46607121  58.91490006  43.59891928  40.18218301]
 [ 47.74709207  56.95186898  44.29555797  51.85307342]
 [ 49.07403549  45.20336146  55.72672011  41.5389062 ]
 [ 42.61495981  56.03420662  41.15230121  54.2535381 ]
 [ 49.45825967  41.66795569  54.66920575  46.86418535]
 [ 39.43105308  55.88452496  48.3999337  37.99631793]
 [ 47.40390326  49.8714146  39.85628813  50.95375799]
 [ 43.13473547  38.23482222  50.53848857  39.05458419]
 [ 40.04532618  52.15449568  51.45401734  55.46855359]
 [ 52.10264408  55.09211195  57.97925338  49.72843111]
 [ 54.85096488  54.57670054  54.08504454  60.85481074]
 [ 46.15207206  51.36887118  62.71209341  52.83290343]
 [ 50.09891239  52.09971137  48.0198356  54.93280663]
 [ 6.          31.385475  60.68224211  43.27978035]]

```

```

Learned Q-values (Q-learning):
[[ 84.          84.          83.          90.          ]
 [ 83.99323558  82.26242627  90.          79.87679232]
 [ 82.99956675  90.          79.9755778  61.75589322]
 [ 90.          80.          62.          79.          ]
 [ 80.          62.          79.          62.          ]
 [ 62.          79.          62.          80.          ]
 [ 79.          62.          80.          77.08160155]
 [ 61.99996635  80.          77.09746188  63.99242975]
 [ 80.          77.107  64.007  75.007  ]
 [ 77.08076925  64.00119391  75.00605486  81.23  ]
 [ 63.99223825  74.99248434  81.22470689  82.23  ]
 [ 75.00387854  81.22561366  82.23  77.6  ]
 [ 81.22966136  82.22993392  77.6  84.7  ]
 [ 82.22999972  77.6  84.7  84.  ]
 [ 77.59999975  69.3273562  84.  83.99999713]
 [ 76.11008349  70.56  50.4  83.  ]]

```

```

Final Policy (SARSA):
[3 2 1 1 2 1 2 1 3 2 3 2 3 2 3 2]

```

```

Final Policy (Q-learning):
[3 2 1 0 0 3 2 1 0 3 3 2 3 2 2 3]

```

```

import numpy as np
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 5))


plt.subplot(1, 2, 1)
plt.imshow(learned_Q_values_sarsa, cmap='viridis')
plt.colorbar()
plt.title('Learned Q-values (SARSA)')
plt.xlabel('Actions')
plt.ylabel('States')

plt.subplot(1, 2, 2)
plt.imshow(learned_Q_values_qlearning, cmap='viridis')
plt.colorbar()
plt.title('Learned Q-values (Q-learning)')
plt.xlabel('Actions')
plt.ylabel('States')


plt.tight_layout()
plt.show()

```

Learned Q-values (SARSA)



Learned Q-values (Q-learning)

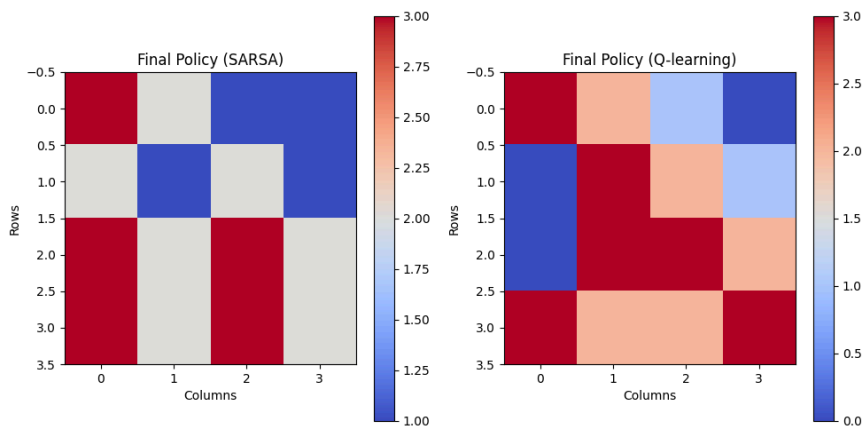


```
plt.figure(figsize=(10, 5))

plt.subplot(1, 2, 1)
plt.imshow(final_policy_sarsa.reshape(num_rows, num_cols), cmap='coolwarm', interpolation='nearest')
plt.colorbar()
plt.title('Final Policy (SARSA)')
plt.xlabel('Columns')
plt.ylabel('Rows')

plt.subplot(1, 2, 2)
plt.imshow(final_policy_qlearning.reshape(num_rows, num_cols), cmap='coolwarm', interpolation='nearest')
plt.colorbar()
plt.title('Final Policy (Q-learning)')
plt.xlabel('Columns')
plt.ylabel('Rows')

plt.tight_layout()
plt.show()
```



```
import numpy as np
import matplotlib.pyplot as plt

combined_Q_values = np.concatenate((learned_Q_values_sarsa, learned_Q_values_qlearning), axis=0)
```

```
plt.figure(figsize=(10, 5))

plt.subplot(1, 2, 1)
```

```

for action in range(num_actions):
    plt.hist(learned_Q_values_sarsa[:, action], bins=20, alpha=0.5, label=f'Action {action}')
plt.xlabel('Q-values')
plt.ylabel('Frequency')
plt.title('Distribution of Learned Q-values (SARSA)')
plt.legend()

plt.subplot(1, 2, 2)
for action in range(num_actions):
    plt.hist(learned_Q_values_qlearning[:, action], bins=20, alpha=0.5, label=f'Action {action}')
plt.xlabel('Q-values')
plt.ylabel('Frequency')
plt.title('Distribution of Learned Q-values (Q-learning)')
plt.legend()

plt.tight_layout()
plt.show()

```

